

Domain-Adversarial Training of Neural Networks

Yaroslav Ganin

GANIN@SKOLTECH.RU

Evgeniya Ustinova

EVGENIYA.USTINOVA@SKOLTECH.RU

Skolkovo Institute of Science and Technology (Skoltech)

Skolkovo, Moscow Region, Russia

Hana Ajakan

HANA.AJAKAN.1@ULaval.CA

Pascal Germain

PASCAL.GERMAIN@IFT.ULaval.CA

Département d'informatique et de génie logiciel, Université Laval

Québec, Canada, G1V 0A6

Hugo Larochelle

HUGO.LAROCHELLE@USHERBROOKE.CA

Département d'informatique, Université de Sherbrooke

Québec, Canada, J1K 2R1

François Laviolette

FRANCOIS.LAVIOLETTE@IFT.ULaval.CA

Mario Marchand

MARIO.MARCHAND@IFT.ULaval.CA

Département d'informatique et de génie logiciel, Université Laval

Québec, Canada, G1V 0A6

Victor Lempitsky

LEMPITSKY@SKOLTECH.RU

Skolkovo Institute of Science and Technology (Skoltech)

Skolkovo, Moscow Region, Russia

Contents

1. Basic concepts




- Transfer learning
- Domain Adaptation
- GAN
- H-Divergence, VC-dimension

2. Method

3. Result

Basic concepts

- Transfer learning



Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]

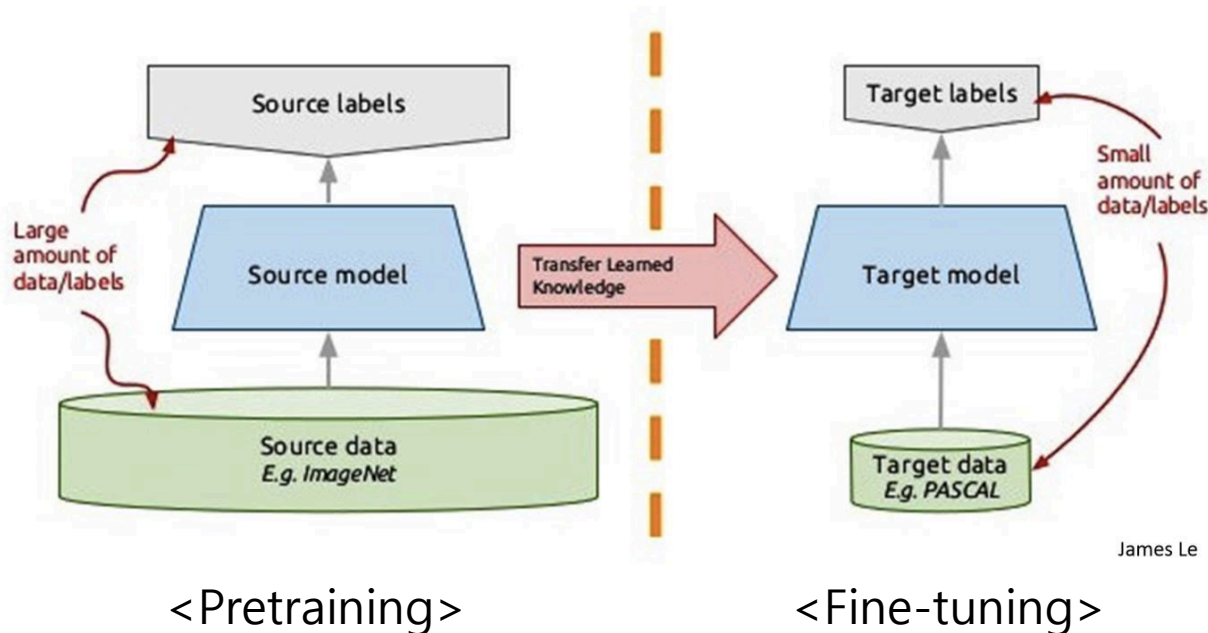
CNN에서의 학습

- CNN을 통하여 Feature (Weight) 를 학습한다.
- CNN layer의 앞쪽은 Low-level feature 를 학습
Low-level feature : Edge, Corner, Color, ...
Object의 요소가 될 수 있는 Feature
- 뒤쪽은 High-level feature 를 학습
High-level feature : Object, Object의 일부분
구체적인 Object에 관하여

Basic concepts

- Transfer learning

Transfer learning: idea

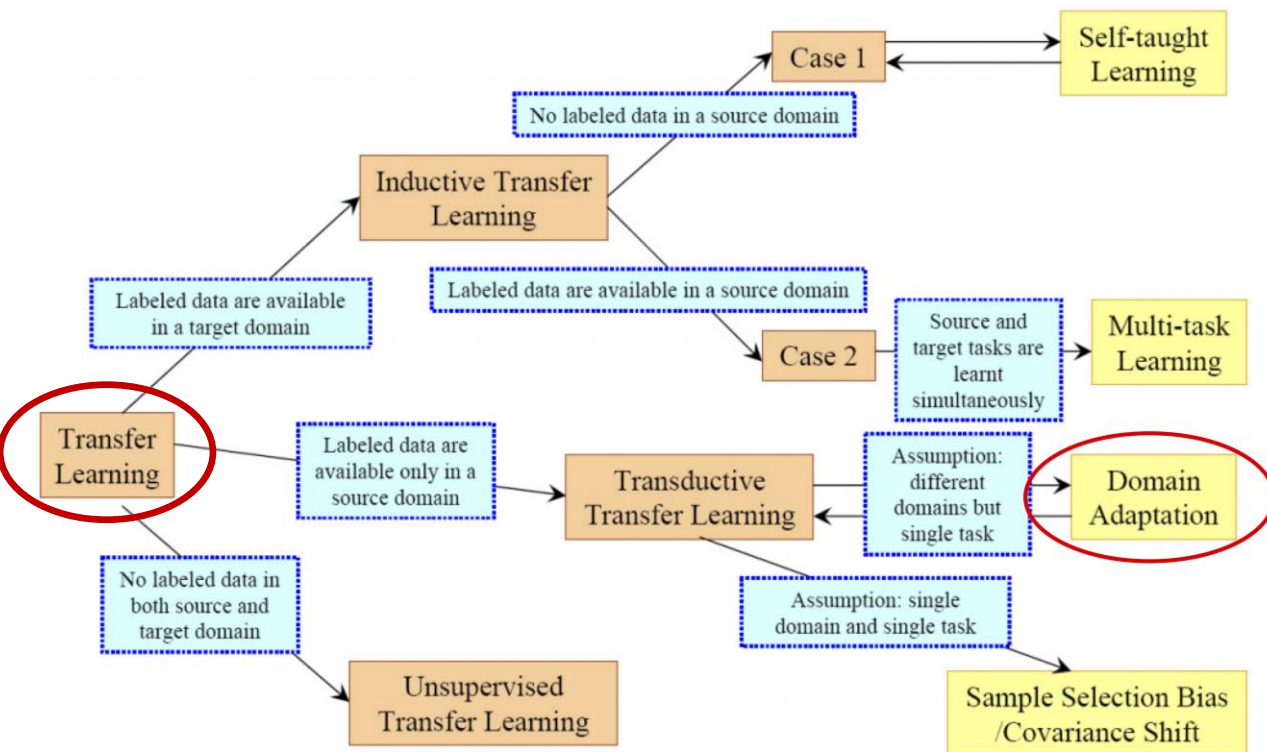


Transfer learning

- 이미 학습된 Network의 Knowledge를 다른 작업에 사용 하는 행위
- 많은 양의 데이터를 가진 Source domain을 학습하는 단계인 'Pretraining'
- 상대적으로 적은 data가진 Target domain을 학습하는 단계인 'Fine tuning'

Basic concepts

- Transfer learning



Transfer learning

Transfer learning은 특정 Task에 학습된 Network의 Knowledge를 다른 작업에 사용 하는 행위를 말한다.

Domain Adaptation

- Transfer learning의 한 종류
- No Target data's label
- Source <-> Target 은 다른 domain에 존재
- Single task, Source와 Target은 같은 task

Basic concepts

- Domain adaptation

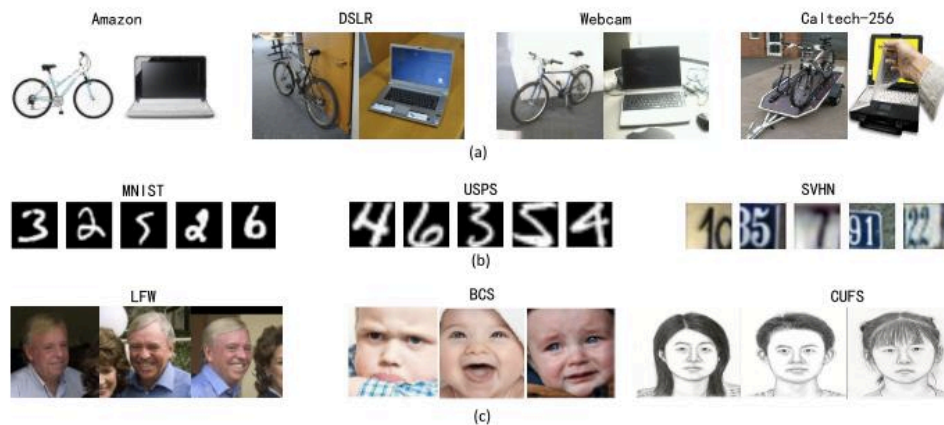
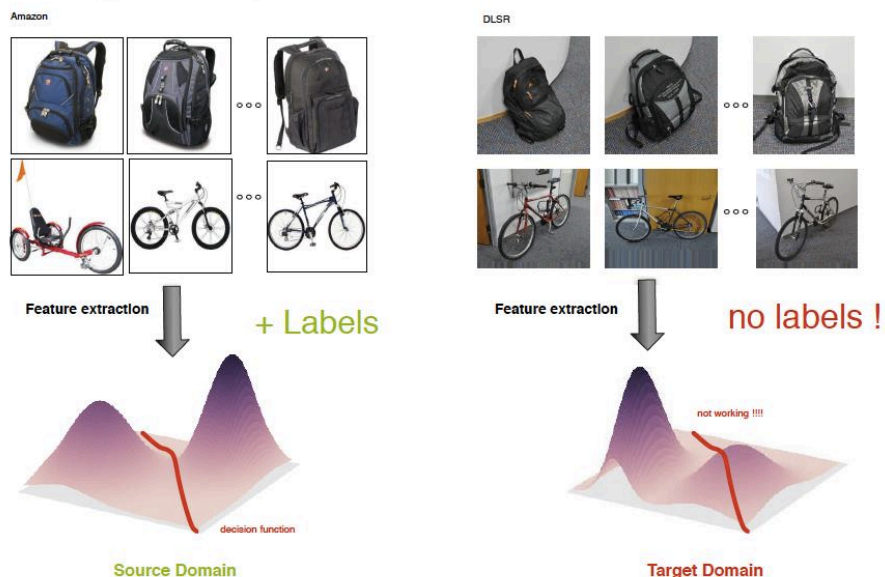


Fig. 1. (a) Some object images from the "Bike" and "Laptop" categories in Amazon, DSLR, Webcam, and Caltech-256 databases. (b) Some digit images from MNIST, USPS, and SVHN databases. (c) Some face images from LFW, BCS and CUFS databases. Realworld computer vision applications, such as face recognition, must learn to adapt to distributions specific to each domain.



Why we need domain adaptation?

1. Domain Shift

- Fig 1. 4개의 Dataset 존재 (맨위)

Amazon, DSLR, Webcam, Caltech, ...

- 각 Dataset은 같은 label을 가지고 있다.

Bicycle, Laptop, ...

- 이중 하나의 Dataset으로 Train 된 classifier가 존재


- Classifier는 다른 Dataset에서도 효과적일까?

- 대부분의 경우 Source에서 만큼의 성능 X

- label이 같더라도 Data의 Distribution이 다르기 때문

Basic concepts

- Domain adaptation



Basic concepts


- Domain adaptation



Why we need domain adaptation?

2. hard to labeling

- Labeling 은 굉장히 귀찮고 시간이 많이 드는 작업
- 전문가의 도움을 받아야 하는 domain들도 다수 존재
- Label된 data가 많으면 좋겠지만
Real-world에서는 그런 경우가 매우 드물다.
Ex. 병원을 처음 오는 환자, Self-driving car 초행길



Basic concepts

- Domain adaptation

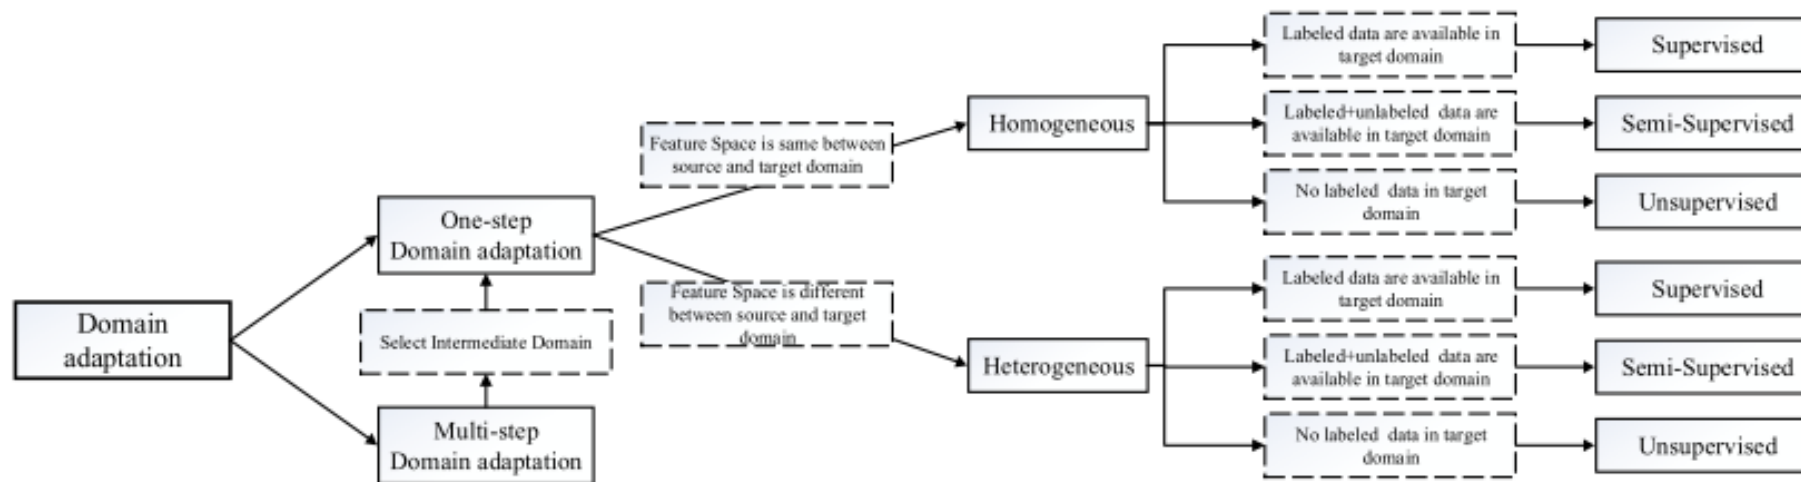


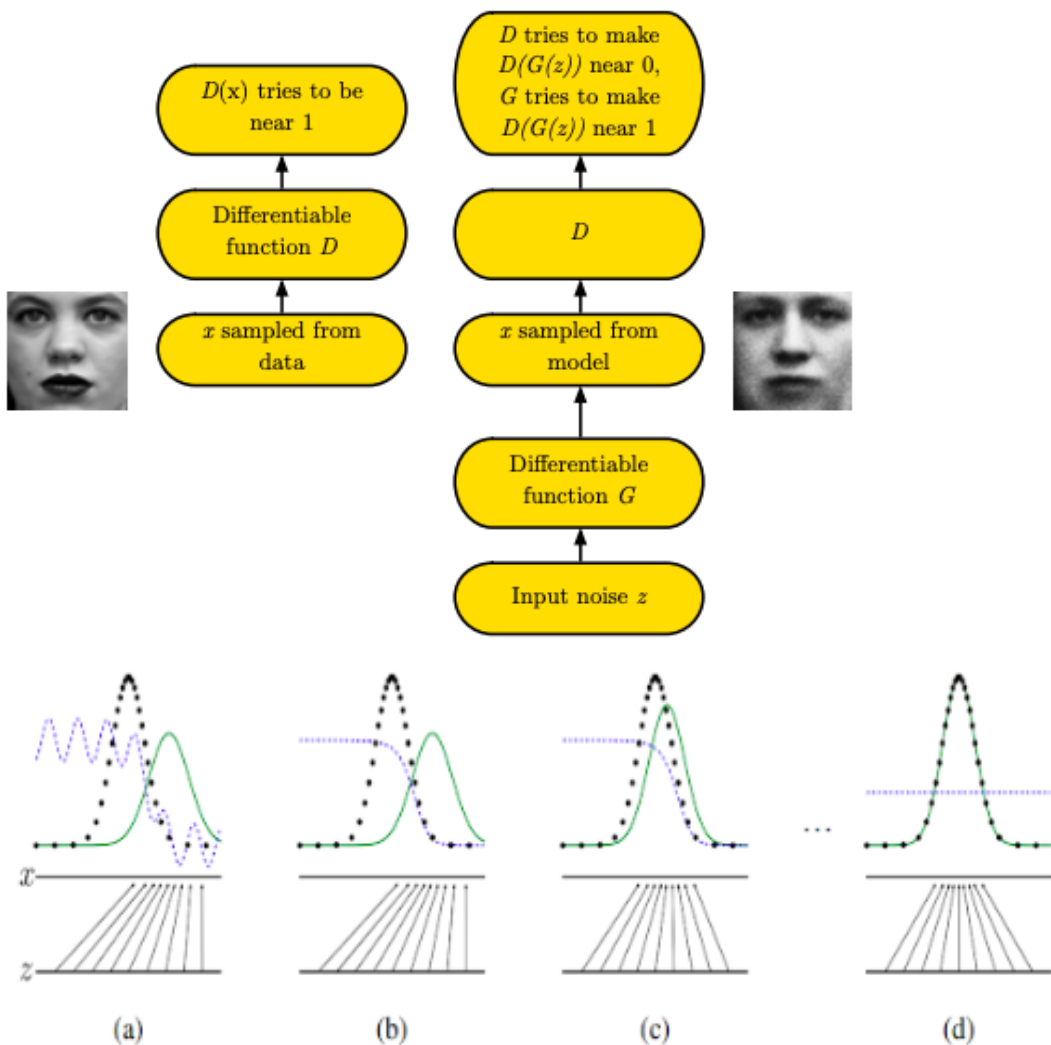
Fig. 2. An overview of different settings of domain adaptation

TABLE I
DIFFERENT DEEP APPROACHES TO ONE-STEP DA

One-step DA Approaches	Brief Description	Subsettings
Discrepancy-based	fine-tuning the deep network with labeled or unlabeled target data to diminish the domain shift	class criterion [118], [86], [79], [98] [53], [45], [75], [139], [130], [29], [118], [28] statistic criterion [74], [130], [73] [75], [120], [32], [109], [87], [144] architecture criterion [69], [54], [68], [95], [128], [89] geometric criterion [16]
Adversarial-based	using domain discriminators to encourage domain confusion through an adversarial objective	generative models [70], [4], [57] non-generative models [119], [118], [26], [25], [117] [85]
Reconstruction-based	using the data reconstruction as an auxiliary task to ensure feature invariance	encoder-decoder reconstruction [5], [33], [31], [144] adversarial reconstruction [131], [143], [59]

Basic concepts

- GAN



- Generative model 'G' 와 Discriminator model 'D' 존재
- 'G'는 data의 distribution을 학습하여 유사하게 생성
- 'D'는 'G'가 생성한 distribution과 실제 data의 distribution을 구별해 내는 역할
- 따라서 'G'는 'D'가 구별해내지 못하도록 distribution 학습
'D'는 잘 구분하도록 학습
- 두번째 그림의 검은 점선이 Data distribution
파란 점선이 'D', 녹색 실선이 'G' 이다.
- 파란 점선은 녹색 실선과 검은 점선을 구별해 내려 한다
- 하지만 'G'는 'D'가 검은 점선과 녹색 실선을 구별하기 힘들게 학습을 한다
- 결국에는 (d)와 같은 형태가 된다.

Basic concepts

- GAN

Kullback-Leibler Divergence

$$D_{KL}(p||q) = E[\log(p_i) - \log(q_i)] = \sum_i p_i \log \frac{p_i}{q_i}$$

Non - Symmetric

$$D_{KL}(p||q) \neq D_{KL}(q||p)$$

Jensen-Shannon Divergence

$$JSD(p, q) = \frac{1}{2} D_{KL}(p||\frac{p+q}{2}) + \frac{1}{2} D_{KL}(q||\frac{p+q}{2})$$

데이터 분포 간의 유사도

Divergence는 두 데이터의 확률 분포의 차이를 나타낸다.

Kullback-Leibler Divergence (KL Divergence)

두 모델(두 데이터)의 확률 분포가 얼마나 유사한지 알 수 있는 척도 값이 적을수록 두 확률 분포가 유사하다.

- Jensen-Shannon Divergence

KL Divergence를 보완한 Divergence

KL Divergence 는 Symmetric 하지 않는데, 이를 보완하여 Symmetric한 꼴로 만듦.

두 확률 분포 사이의 Distance 역할을 한다.

Basic concepts

- Basic Notation

We consider classification tasks where X is the input space and $Y = \{0, 1, \dots, L-1\}$ is the set of L possible labels. Moreover, we have two different distributions over $X \times Y$, called the *source domain* \mathcal{D}_S and the *target domain* \mathcal{D}_T . An *unsupervised domain adaptation* learning algorithm is then provided with a *labeled source sample* S drawn *i.i.d.* from \mathcal{D}_S , and an *unlabeled target sample* T drawn *i.i.d.* from \mathcal{D}_T^X , where \mathcal{D}_T^X is the marginal distribution of \mathcal{D}_T over X .

$$S = \{(\mathbf{x}_i, y_i)\}_{i=1}^n \sim (\mathcal{D}_S)^n; \quad T = \{\mathbf{x}_i\}_{i=n+1}^N \sim (\mathcal{D}_T^X)^{n'},$$

with $N = n + n'$ being the total number of samples. The goal of the learning algorithm is to build a classifier $\eta : X \rightarrow Y$ with a low *target risk*

$$R_{\mathcal{D}_T}(\eta) = \Pr_{(\mathbf{x}, y) \sim \mathcal{D}_T} (\eta(\mathbf{x}) \neq y),$$

while having no information about the labels of \mathcal{D}_T .



Figure 6: Examples of domain pairs used in the experiments. See Section 5.2.4 for details.

- X 는 Input space (Dataset)
- Y 는 Label 집합
- $\mathcal{D}_{\{S\}}, \mathcal{D}_{\{T\}}$ 는 각 Domain 을 의미한다.
- η 는 Classifier를 의미한다.
- $R_{\mathcal{D}_T}(\eta)$ 는 Classifier의 empirical error

Basic concepts

- \mathcal{H} -divergence / VC-dimension

Definition 1 (Ben-David et al., 2006, 2010; Kifer et al., 2004) *Given two domain distributions \mathcal{D}_S^X and \mathcal{D}_T^X over X , and a hypothesis class \mathcal{H} , the \mathcal{H} -divergence between \mathcal{D}_S^X and \mathcal{D}_T^X is*

$$d_{\mathcal{H}}(\mathcal{D}_S^X, \mathcal{D}_T^X) = 2 \sup_{\eta \in \mathcal{H}} \left| \Pr_{\mathbf{x} \sim \mathcal{D}_S^X} [\eta(\mathbf{x}) = 1] - \Pr_{\mathbf{x} \sim \mathcal{D}_T^X} [\eta(\mathbf{x}) = 1] \right|.$$

That is, the \mathcal{H} -divergence relies on the capacity of the hypothesis class \mathcal{H} to distinguish between examples generated by \mathcal{D}_S^X from examples generated by \mathcal{D}_T^X . Ben-David et al. (2006, 2010) proved that, for a symmetric hypothesis class \mathcal{H} , one can compute the *empirical \mathcal{H} -divergence* between two samples $S \sim (\mathcal{D}_S^X)^n$ and $T \sim (\mathcal{D}_T^X)^{n'}$ by computing

$$\hat{d}_{\mathcal{H}}(S, T) = 2 \left(1 - \min_{\eta \in \mathcal{H}} \left[\frac{1}{n} \sum_{i=1}^n I[\eta(\mathbf{x}_i) = 0] + \frac{1}{n'} \sum_{i=n+1}^N I[\eta(\mathbf{x}_i) = 1] \right] \right), \quad (1)$$

where $I[a]$ is the indicator function which is 1 if predicate a is true, and 0 otherwise.

Theorem 2 (Ben-David et al., 2006) *Let \mathcal{H} be a hypothesis class of VC dimension d . With probability $1 - \delta$ over the choice of samples $S \sim (\mathcal{D}_S^X)^n$ and $T \sim (\mathcal{D}_T^X)^n$, for every $\eta \in \mathcal{H}$:*


$$R_{\mathcal{D}_T}(\eta) \leq R_S(\eta) + \sqrt{\frac{4}{n} \left(d \log \frac{2en}{d} + \log \frac{4}{\delta} \right)} + \hat{d}_{\mathcal{H}}(S, T) + 4 \sqrt{\frac{1}{n} \left(d \log \frac{2n}{d} + \log \frac{4}{\delta} \right)} + \beta,$$

with $\beta \geq \inf_{\eta^* \in \mathcal{H}} [R_{\mathcal{D}_S}(\eta^*) + R_{\mathcal{D}_T}(\eta^*)]$, and

$$R_S(\eta) = \frac{1}{n} \sum_{i=1}^n I[\eta(\mathbf{x}_i) \neq y_i]$$


Basic concepts

- H -divergence




Instances X

Hypotheses H




$x_1 = \langle \text{Sunny Warm Normal Strong Warm Same} \rangle, +$
 $x_2 = \langle \text{Sunny Warm High Strong Warm Same} \rangle, +$
 $x_3 = \langle \text{Rainy Cold High Strong Warm Change} \rangle, -$
 $x_4 = \langle \text{Sunny Warm High Strong Cool Change} \rangle, +$

$h_0 = \langle \emptyset, \emptyset, \emptyset, \emptyset, \emptyset, \emptyset \rangle$
 $h_1 = \langle \text{Sunny Warm Normal Strong Warm Same} \rangle$
 $h_2 = \langle \text{Sunny Warm ? Strong Warm Same} \rangle$
 $h_3 = \langle \text{Sunny Warm ? Strong Warm Same} \rangle$
 $h_4 = \langle \text{Sunny Warm ? Strong ? ?} \rangle$



Computed by WolframAlpha



Computed by WolframAlpha

- Hypothesis space는 Classifier가 존재하는 공간
- Model complexity라 생각하면 됨
ex) Linear regression -> limited to linear functions as its hypothesis space
- 일종의 Hyperparameter

Basic concepts

- H -divergence


$$d_{\mathcal{H}}(\mathcal{D}_S^X, \mathcal{D}_T^X) = 2 \sup_{\eta \in \mathcal{H}} \left| \Pr_{\mathbf{x} \sim \mathcal{D}_S^X} [\eta(\mathbf{x}) = 1] - \Pr_{\mathbf{x} \sim \mathcal{D}_T^X} [\eta(\mathbf{x}) = 1] \right|.$$
$$\hat{d}_{\mathcal{H}}(S, T) = 2 \left(1 - \min_{\eta \in \mathcal{H}} \left[\frac{1}{n} \sum_{i=1}^n I[\eta(\mathbf{x}_i) = 0] + \frac{1}{n'} \sum_{i=n+1}^N I[\eta(\mathbf{x}_i) = 1] \right] \right),$$

- Divergence -> data distribution을 parameter로

- H space(η 가 존재하는)에 영향을 받는다.

- H 의 값이 커지면 두 distribution을 잘 구분

- 따라서 H -divergence 값도 커진다.



Basic concepts


- VC Dimension

- **Example:** VC-dimension of **spherical indicator functions**.

Consider spherical decision surfaces in a d -dimensional \mathbf{x} -space, parameterized by center \mathbf{c} and radius r parameters:

$$f(\mathbf{x}, \mathbf{c}, r) = I((\mathbf{x} - \mathbf{c})^2 \leq r^2)$$

In a 2-dim space ($d=2$) there exists 3 points that can be shattered, but 4 points cannot be shattered $\rightarrow h=3$



- Classification algorithm의 capacity 측정
- $VC(H)$ 란 현재 H -space에서 label에 상관없이 정확하게 분류할 수 있는 가장 큰 X (Data)
- 왼쪽 그림의 경우 $H = 3$ 이다.

<http://sanghyukchun.github.io/66/> 참조

Basic concepts

Theorem 2 (Ben-David et al., 2006) Let \mathcal{H} be a hypothesis class of VC dimension d . With probability $1 - \delta$ over the choice of samples $S \sim (\mathcal{D}_S)^n$ and $T \sim (\mathcal{D}_T^x)^n$, for every $\eta \in \mathcal{H}$:

$$\text{Minimize } R_{\mathcal{D}_T}(\eta) \leq R_S(\eta) + \sqrt{\frac{4}{n} \left(d \log \frac{2en}{d} + \log \frac{4}{\delta} \right)} + \hat{d}_{\mathcal{H}}(S, T) + 4\sqrt{\frac{1}{n} \left(d \log \frac{2n}{d} + \log \frac{4}{\delta} \right)} + \beta,$$

with $\beta \geq \inf_{\eta^* \in \mathcal{H}} [R_{\mathcal{D}_S}(\eta^*) + R_{\mathcal{D}_T}(\eta^*)]$, and

$$R_S(\eta) = \frac{1}{n} \sum_{i=1}^m I[\eta(\mathbf{x}_i) \neq y_i]$$

Method






Figure 1: The **proposed architecture** includes a deep *feature extractor* (green) and a deep *label predictor* (blue), which together form a **standard feed-forward architecture**. **Unsupervised domain adaptation** is achieved by adding a *domain classifier* (red) connected to the feature extractor via a *gradient reversal layer* that **multiplies** the gradient by a certain **negative constant** during the **backpropagation-based training**. **Otherwise**, the training proceeds **standardly** and **minimizes** the label prediction **loss** (for **source examples**) and the domain **classification loss** (for **all samples**). **Gradient reversal ensures** that the feature **distributions** over the two domains are made **similar** (as **indistinguishable as possible** for the domain **classifier**), **thus resulting** in the domain-invariant **features**.


Method



(a) MNIST architecture; inspired by the classical LeNet-5 (LeCun et al., 1998).



(b) SVHN architecture; adopted from Srivastava et al. (2014).



(c) GTSRB architecture; we used the single-CNN baseline from Cireşan et al. (2012) as our starting point.

Figure 4: CNN architectures used in the experiments. Boxes correspond to transformations applied to the data. Color-coding is the same as in Figure 1.

Method

```
def __init__(self):
    super(CNNModel, self).__init__()
    self.feature = nn.Sequential()
    self.feature.add_module('f_conv1', nn.Conv2d(3, 64, kernel_size=5))
    self.feature.add_module('f_bn1', nn.BatchNorm2d(64))
    self.feature.add_module('f_pool1', nn.MaxPool2d(2))
    self.feature.add_module('f_relu1', nn.ReLU(True))
    self.feature.add_module('f_conv2', nn.Conv2d(64, 50, kernel_size=5))
    self.feature.add_module('f_bn2', nn.BatchNorm2d(50))
    self.feature.add_module('f_drop1', nn.Dropout2d())
    self.feature.add_module('f_pool2', nn.MaxPool2d(2))
    self.feature.add_module('f_relu2', nn.ReLU(True))

    self.class_classifier = nn.Sequential()
    self.class_classifier.add_module('c_fc1', nn.Linear(50 * 4 * 4, 100))
    self.class_classifier.add_module('c_bn1', nn.BatchNorm1d(100))
    self.class_classifier.add_module('c_relu1', nn.ReLU(True))
    self.class_classifier.add_module('c_drop1', nn.Dropout2d())
    self.class_classifier.add_module('c_fc2', nn.Linear(100, 100))
    self.class_classifier.add_module('c_bn2', nn.BatchNorm1d(100))
    self.class_classifier.add_module('c_relu2', nn.ReLU(True))
    self.class_classifier.add_module('c_fc3', nn.Linear(100, 10))
    self.class_classifier.add_module('c_softmax', nn.LogSoftmax(dim=1))

    self.domain_classifier = nn.Sequential()
    self.domain_classifier.add_module('d_fc1', nn.Linear(50 * 4 * 4, 100))
    self.domain_classifier.add_module('d_bn1', nn.BatchNorm1d(100))
    self.domain_classifier.add_module('d_relu1', nn.ReLU(True))
    self.domain_classifier.add_module('d_fc2', nn.Linear(100, 2))
    self.domain_classifier.add_module('d_softmax', nn.LogSoftmax(dim=1))
```

Result

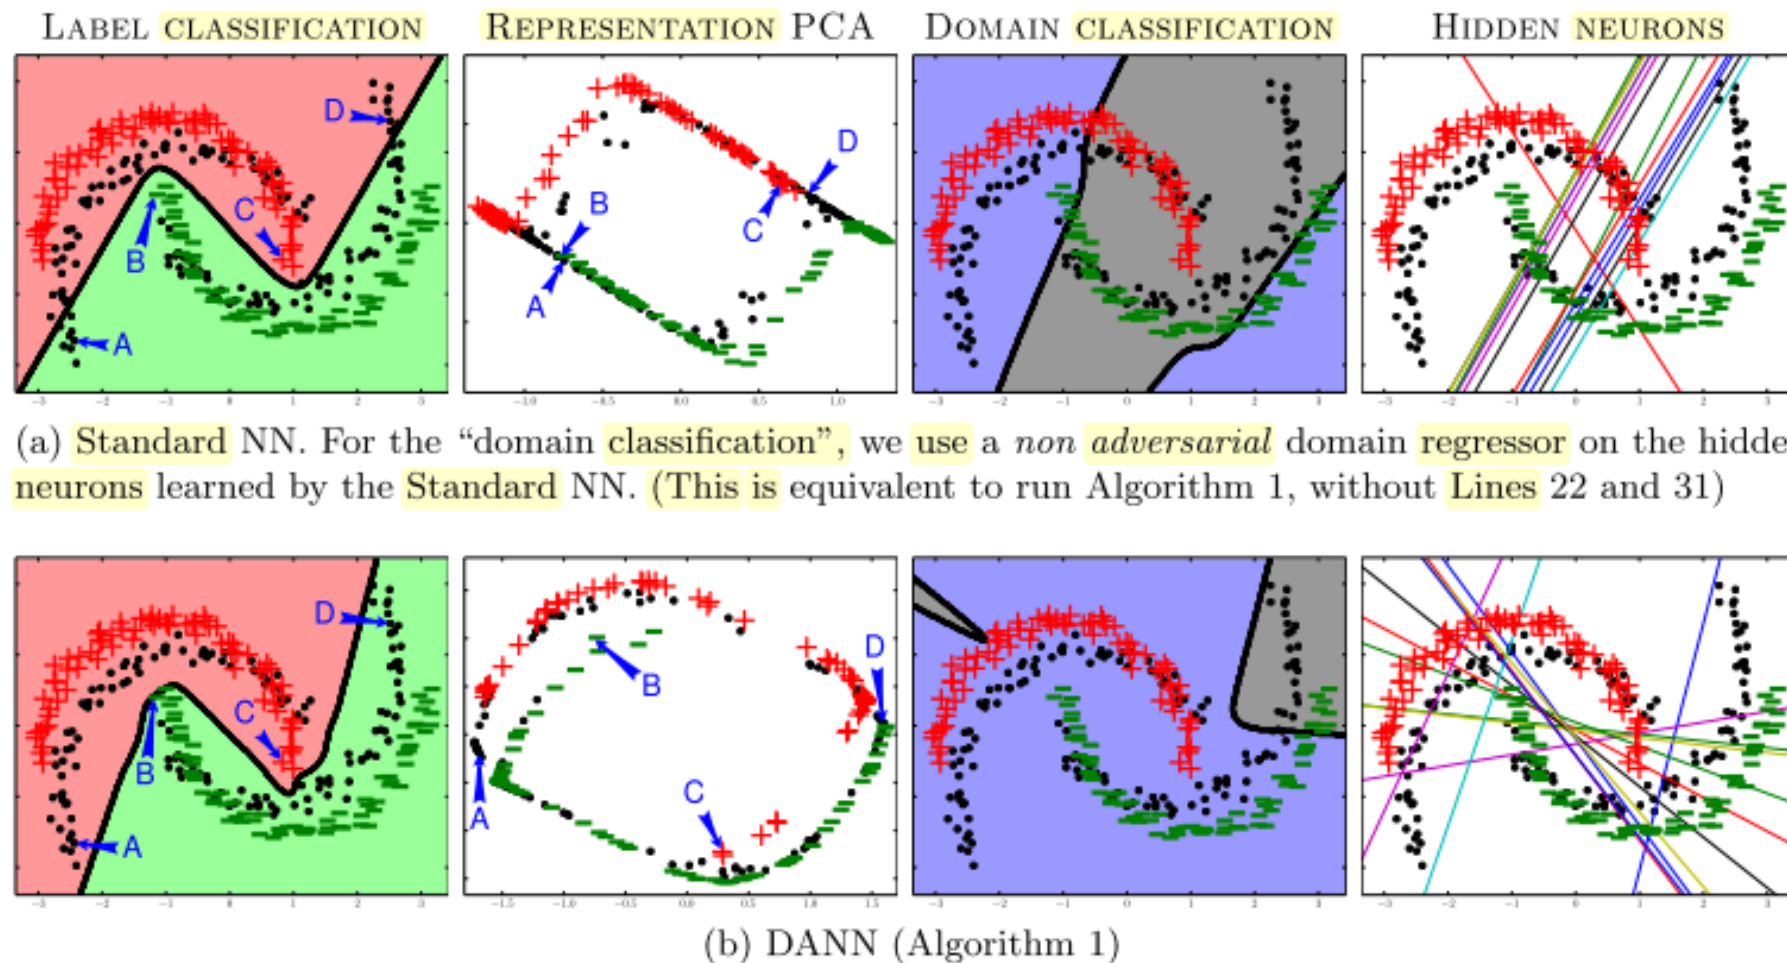


Figure 2: The *inter-twinning moons* toy problem. Examples from the source sample are represented as a “+” (label 1) and a “-” (label 0), while examples from the unlabeled target sample are represented as black dots. See text for the figure discussion.